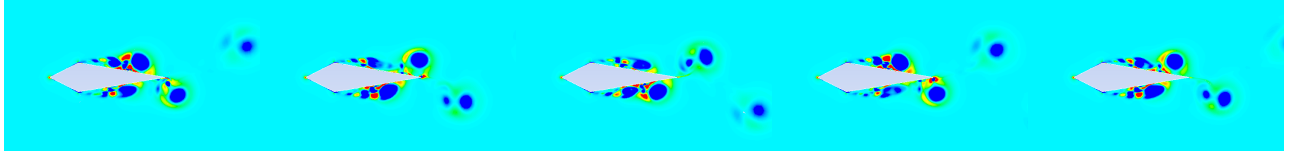


# Real-time Aerodynamic Sound Synthesis for Slender Objects

Jui-Hsien Wang\*  
Cornell University



**Figure 1:** The vortex shedding pattern of flow over a 2D sword determines its dipole characteristics and the aerodynamic sound when swung by, say, a virtual samurai.

## Abstract

In this project, I explored the problem of real-time, physics-based aerodynamic sound synthesis for slender objects. It is largely inspired by the paper by Dobashi *et al.* [Dobashi et al. 2003]. The aerodynamic sound when we swing a slender object, such as a stick, is originated from the complex interaction between the air flow and the stick. Sufficient spatial/temporal resolution was regarded to be essential to capture the physics and thus the characteristics of the sound generated. However, the required fluid simulation is too expensive to run at run-time audio stepping rate. To avoid such computation, I first precomputed a comprehensive database that contains relevant sound textures evaluated from high-quality grid-based fluid simulation, and then at runtime, this database is fetched and textures are blended to effectively resynthesize the aerodynamic swinging sound. Next, to increase the interactivity of the project, I interfaced the sound system with Leap Motion sensor to give real-time motion capture data. The system is proven to be quite reliable and can run at real-time even on a low-end laptop, and create realistic swinging sound.

**Keywords:**

## 1 Introduction

In games and virtual environments, accurately representing the sound of moving a slender object quickly, such as a character swinging a club, is important for story telling and the realism of the environment. However, most of the time, this sound is faked by either playing back experimentally recorded “canned” sound, or by using random parametric models. These models are cheap to evaluate, but they lack physical basis of how the sound is generated, and therefore can cause noticeable audio-visual synchronization problem or need a lot of hand-tuning to ensure the quality. Dobashi [Dobashi et al. 2003] presented the first automated, physically-based aerodynamic sound synthesis pipeline.

\*e-mail: jw969@cornell.edu

This class of aerodynamic sound is generated by the fluid flow around the object, and can be described by Curle’s model [?; ?]. Given two critical assumptions: (1) the object is acoustically compact, and (2) the listener is placed at a far-field position, the aerodynamic sound can be approximated by a dipole source, whose magnitude is governed by the unsteady forces generated by the object placed in the flow field. Section 3 specifies more details on this model.

Using this model, we computed fluid flow around the object of interest and the dipole sound source function, and cached them in a database for that object for runtime evaluation. At runtime, the propagation is modeled by use of free-space Green’s function, and the sound source function was used to evaluate the sound at listener’s position given the motion of object. Amplitude scaling based on Curle’s model and frequency scaling based on Strouhal number were implemented. Section 4 and Section 5 discuss these parts.

Leap Motion was configured to give a run-time motion capture data, and used as the input of our sound system. The basics of this interface will be described and demonstrated in Section 5.

## 2 Related Work

## 3 Curle’s Model

For acoustically compact object and far-field listener, the sound pressure is governed by the equation [?]

$$p(\mathbf{q}, t) = \frac{1}{4\pi c_0 r^2} (\mathbf{q} - \mathbf{o}) \cdot \mathbf{g}(t - \frac{r}{c_0}), \quad (1)$$

$$\mathbf{g}(t) = \frac{\partial}{\partial t} \mathbf{F}(t) dS. \quad (2)$$

$\mathbf{q}$  is the listener position, and  $\mathbf{o}$  is the object center position.  $c_0$  is the speed of sound,  $r$  is the distance between listener and object center, and  $\mathbf{F}$  is the aerodynamic forces of the object in the texture domain, precomputed in the fluid simulation. Given this model, we can statically sample and cache the function  $\mathbf{g}$  at different inflow speed and directions for a given object, and only evaluate the sound propagation at runtime. The runtime cost would be a data-fetching process, and a few flops to decide the position and direction of the object with respect to the listener in order to scale the sound amplitude and time delay.

## 4 Sound Texture Database Construction

The sound source function,  $\mathbf{g}$ , can be precomputed in a fluid simulation. I used proprietary finite-volume solver Ansys Fluent to do

the computation. The object being sampled was placed in a virtual wind tunnel and rotated with 5 degree interval to sample the sound source function for every inflow direction. To shorten this preprocess, the following identities are used

- (1) The amplitude of the dipole sound scales linearly with inflow speed in the log-space. That is,  $|p_v| = (v/v_0)^\alpha |p_{v_0}|$ .  $\alpha = 6$  is the coefficient Dobashi used, but I found that  $\alpha \approx 2$  works better because the high coefficient gives too wide the dynamic range.
- (2) The frequency of the dipole sound scales linearly with inflow speed. The physical argument here is based on the Strouhal frequency,  $St = fD/v$ , where  $St$  is the Strouhal number,  $D$  is the diameter of the cylinder, and  $f$  is the frequency. For a wide range of Reynold's number ( $Re = 100 \sim 1E6$ ), the Strouhal number stays roughly constant for cylinder ( $St_{cylinder} \sim 0.21$ ) and therefore for the same object, the frequency scales linearly with inflow velocity. Linear interpolation was applied to sound textures to first upsample it to avoid artifacts when we scale the textures.
- (3) For inflow with non-zero incidence angle, the sound is approximately equal to the sound created by inflow speed scaled by the geometric factor. Therefore, all the inflow directions at runtime were projected to the normal direction of the object before the sound source function was evaluated.

$$sound(v_\theta) \approx sound(v_\parallel), \quad v_\parallel = v \cos \theta. \quad (3)$$

## 5 Runtime Considerations


### 5.1 Real-time motion capture

Position, speed, and orientation of the object are the metrics needed to compute the aerodynamic sound. To test the system, I first drove the sound model using mouse cursor (only speed was used), and then I configured Leap Motion to obtain the position, speed, and orientation information in real-time. Intertextural blending is introduced at the second stage.

**Using mouse cursor** To drive the sound model using mouse, the cursor position in screen space was queried at every audio callback, and the speed of the cursor was computed using first-order explicit Euler method. No orientation effect was introduced at this stage (therefore only one texture was considered).


**Using Leap Motion** Next, I sought for low latency, stable sensors to compute runtime object orientation. I have tried Nintendo Wii's controller but the bluetooth interface was not stable, and the field-of-view (FOV) of its infrared sensor was too small to resolve swinging motion. I then found Leap Motion, which is a low-cost commercial sensor released in 2013. The mature API makes it quite easy to use. The Leap Motion system recognizes and tracks hands, fingers and finger-like tools. The device operates in an intimate proximity with high precision and tracking frame rate (on my machine the peak sampling rate can reach 150Hz when lighting condition is good) and reports discrete positions, gestures, and motion. The Leap Motion controller uses optical sensors and infrared light. The sensors have FOV of about 150 degrees, making it a good candidate for the swinging application. Please see Fig. 2 for more detailed description of this system.

Knowing the direction of the hand,  $\mathbf{h}$ , and the normal direction of the palm,  $\mathbf{p}$ , the relative wind direction in the world space can be transformed to the object space. Using the identity for the slender



**Figure 2:** Leap Motion sensor: (a) The system use IR sensors to detect the motion and gestures of the hand. The raw IR images can be used for debugging; (b) The system features a 150 degree FOV, making it a good candidate for our swinging motion; (c) The orientation of the object can be uniquely determined by transforming the wind vector into the coordinate system defined by the direction of the hand and the palm normal.

object Eq. 3, the sound source function can then be fetched by projecting this wind direction onto the plane of the cross-section and quantized to an integer number that represents the index of the texture (5 degree interval for all cases). This process is illustrated in Fig. 3.



**Figure 3:** The wind direction, determined by the swinging motion, is first transformed into the object space, and then projected onto the plane of the cross-section. This vector,  $\mathbf{w}_\parallel^o$ , is then quantized into a texture index for this object to fetch the sound database, which is preloaded when the program starts.

### 5.2 Clicking noise and audio underrun

One of the big problems I encountered when implementing this paper was how to do the real-time audio rendering properly, the discussion of which is missing from the paper. There were two main problems associated to it: (1) the time-scale mismatch in motion capture refresh rate and audio stream sampling rate can cause severe “clicking” sound when scaling the textures or jumping between textures. (2) the audio underrun caused by improper thread priority for audio callback and high motion capture latency.

The clicking problem (either caused by resampling the texture in the amplitude/frequency shifting or jumping between textures) can be resolved by requesting an optimum buffer size. To avoid audio glitch, the audio interface (I use portaudio) reduces the number of audio callbacks by requesting an audio buffer when the audio stream is opened. The size of the buffer can be fixed or adaptive, depending on the hardware implementation and the application (for example, the CoreAudio framework in MacOS is well written, and can largely reduce the latency). What I did to resolve the clicking was to enforce the motion-capture data to be sent only when audio callback happens, and then linearly blend between the two states of motion data (sound amplitude/frequency). I found that the buffer size of  $\sim 100$  works the best. If the buffer is too large, the motion-capture data is updated infrequently and the playback will sound sluggish; whereas if its too small, then the buffer doesn't have enough time to blend the textures and clicking will occur. Note that this number is roughly the ratio between audio sampling rate (for

most of them is 10000Hz) and motion-capture update refresh rate (at around 100Hz).

The audio underrun happens when the buffer is not properly filled in the requested time frame, and portaudio has no choice but to substitute it with zeros. This problem can be severe if the implementation is not handling it explicitly. By requesting high-priority threads to the audio callback and to the motion-capture callback, the underrun problem can be alleviated if the space of sound source function for the object has reasonable size. Also, unbounded time operations such as file I/O should be completely avoided or minimized in the audio callback. For the same reason, the sound textures should all be loaded when the program starts, to avoid unnecessary underrun at the cost of higher memory footprint.

## 6 Result

### 6.1 Implementation Details

The project was implemented in C++. PortAudio<sup>1</sup> was used for audio control. Leap Motion API<sup>2</sup> was used to control the device and interface with the sound system developed based on Dobashi's model. All the fluid simulation data was obtained using Ansys Fluent, a proprietary finite-volume solver. The simulation ran on 32 core desktop equipped with Intel Xeon processors. More statistics of the precomputation is given in Table 1.

### 6.2 2D cylinder

### 6.3 2D square

### 6.4 2D, anisotropic sword

## Acknowledgements

## References

DOBASHI, Y., YAMAMOTO, T., AND NISHITA, T. 2003. Real-time rendering of aerodynamic sound using sound textures based on computational fluid dynamics. *ACM Trans. on Graphics (Proc. SIGGRAPH)* 22, 3.

---

<sup>1</sup><http://www.portaudio.com/>

<sup>2</sup><https://developer.leapmotion.com>

**Table 1:** *Precomputation statistics.*

Shape	Dimension	Sampling resolution (deg)	# Orientations	Cell number	Processing time / sample (mins)
Cylinder	2	5	1	18816	104.27
Square	2	5	10	42308	154.51
Sword	2	5	19	42308	123.46
Anisotropic Sword	2	5	37	31920	101.21